

Speaker Cards for presentation_BEEN

These are short cue cards for the current deck in `presentation_BEEN.tex`. They are written for natural delivery, not as a word-for-word script.

Main Deck

1. Title

Target time: 15-20 sec

This talk asks a simple question: if making a recommender system more politically aligned requires costly user input, will people actually provide that input or will they free-ride on others? The larger claim is that political alignment is not just a design problem for engineers or platforms. It is also a collective-action problem among users of a shared system.

Main point to land: political alignment can be underprovided even when many users want more of it.

2. Political alignment in recommender systems

Target time: 30-40 sec

Recommender systems shape attention, exposure, and sometimes even political or social learning. But once we move beyond pure accuracy, it becomes much less obvious what the system should be aligned to. The key question is no longer only how to recommend, but what counts as the right content and whose values define that standard.

Main point to land: the alignment problem is normative before it is technical.

3. Political alignment as a public good

Target time: 35-45 sec

Improving a recommender often requires costly input from users: ratings, feedback, and value-relevant information. Once the system improves, however, the benefits are shared. That means users may want a better-aligned system while still preferring that someone else bears the cost of improving it.

Main point to land: alignment has the structure of a public good, so free-riding is a central risk.

4. Contribution & Research Questions

Target time: 5-10 sec

Quick signpost only. I use this slide to mark the move from motivation to the paper's contribution and the empirical questions.

Main point to land: we now move from problem statement to contribution.

5. Objective and contribution

Target time: 40-50 sec

The objective is to understand when users are willing to make costly contributions to improve a shared recommender under value alignment or misalignment. The contribution has three parts. First, we reframe political alignment as a public-good dilemma. Second, we use a light economic model to derive behavioral predictions. Third, we test those predictions in an experimental design that maps closely onto the theory.

Main point to land: the paper contributes a framing, a model, and an experiment.

6. Research questions

Target time: 30-40 sec

The first question is whether perceived political alignment affects willingness to contribute to a shared recommender. The second is whether free-riding is stronger with aligned partners or misaligned partners. The third is whether these effects differ between political and non-political domains.

Main point to land: the paper asks how similarity, misalignment, and domain shape contribution behavior.

7. Theoretical Framework

Target time: 5-10 sec

Quick signpost only. I use this slide to tell the audience that the next few slides give the mechanism behind the hypotheses.

Main point to land: the next section is intuition first, not heavy algebra.

8. Model intuition

Target time: 40-50 sec

Start from a shared recommender used by people with different political preferences. Users engage more when the system is closer to their ideal political level, but moving the system in that direction requires costly contributions. Because the same system serves multiple users, part of the benefit from any one person's contribution spills over to others.

Main point to land: costly influence over a shared system creates strategic interdependence.

9. A simple model

Target time: 40-50 sec

Two user groups choose contributions, and the platform then selects the overall political level of the recommender. Higher contributions make a higher political level more attractive for the platform, but contributing is costly for users. So each group wants the system to move closer to its own ideal while hoping others help pay for that shift.

Main point to land: alignment is modeled as a collective-action problem around a common platform choice.

10. Equilibrium intuition

Target time: 45-55 sec

I would not dwell on the equation. The verbal takeaway is that contribution incentives must balance across user types. Users who are more sensitive to political misalignment have more reason to contribute, but because both users affect the same system, each also has a reason to lean on the other's effort.

Main point to land: the equilibrium condition formalizes why willingness to contribute depends on sensitivity to misalignment.

11. Main behavioral implication

Target time: 30-40 sec

This is the payoff-relevant implication. Contributions are strategic substitutes. If I expect a similar partner to contribute more, my own best response may be to contribute less. That is the core logic behind the free-riding prediction.

Main point to land: alignment can foster free-riding precisely because partners want similar outcomes.

12. Experimental Design

Target time: 5-10 sec

Quick signpost only. I use this slide to transition from theory to the behavioral implementation.

Main point to land: the next section shows how the theory is translated into an experiment.

13. Experimental design: overview

Target time: 35-45 sec

The design has three phases. First, a survey and preference elicitation stage gives us the matching information and a ground-truth movie ranking. Second, participants play a ten-round contribution game with stranger matching. Third, a transparent prediction-and-payment stage maps total contributions into recommendation accuracy and then into payoffs.

Main point to land: the design cleanly separates preferences, contribution behavior, and realized incentives.

14. Phase 1: Pre-treatment survey & ground truth

Target time: 40-50 sec

In the first phase, we elicit both political positions and non-political preferences to classify pairs as similar or opposed. We also ask each participant to rank five movies from a common pool. That ranking becomes the ground truth used later to evaluate recommendation accuracy and determine part of the payment.

Main point to land: phase 1 gives us both the treatment labels and the accuracy benchmark.

15. Phase 2: Rating rounds

Target time: 35-45 sec

Participants then enter ten rounds with stranger matching. In each round they see one movie, the cost of rating, and whether the partner is similar or opposed. Importantly, there is no chat, no identity, and no partner history. The learning rule stays fixed; only the matching context changes across treatments.

Main point to land: the strategic environment changes through context, not through the algorithm.

16. Phase 2: The contribution decision

Target time: 30-40 sec

The core choice is binary and simple: rate or skip. Rating costs money now, but it improves the recommender for both members of the pair. That is exactly the social dilemma we want: a privately costly action with shared benefits.

Main point to land: the experiment operationalizes contribution as a clean rate-versus-skip choice.

17. Phase 2: Movie stimuli

Target time: 30-40 sec

We use fictional movies to reduce familiarity effects and strong priors tied to real titles. A separate pre-test selected ten movies from an initial pool of sixteen,

balanced across political and non-political content and also across genres. So the stimulus set is standardized and validated before entering the main task.

Main point to land: the movie pool is designed to carry content variation without relying on prior brand recognition.

18. Phase 3: Individual payoff

Target time: 40-50 sec

The payoff formula has three moving parts: the fixed endowment, the number of ratings a participant submits, and the bonus from how accurately the recommender reconstructs the participant's ground-truth ranking. So the participant trades off a sure contribution cost now against a probabilistic improvement in later accuracy and bonus payments.

Main point to land: the incentives make contribution costly but potentially worthwhile through better recommendation accuracy.

19. Phase 3: From ratings to accuracy

Target time: 35-45 sec

This slide makes the learning rule transparent. Accuracy depends on total ratings from both members of the pair. No ratings gives about one correct position on average; full mutual rating moves that to about three. The key feature is that one person's ratings improve the shared learning environment.

Main point to land: total joint contribution, not just my own effort, determines the system's learning quality.

20. Phase 3: Incentive structure

Target time: 35-45 sec

Now the public-good logic becomes explicit. A rating costs the contributor twenty-five cents, but it increases expected bonus by about twenty cents for each participant. So the private return is below the private cost, while the joint return exceeds the cost. That is why this is a Prisoner's-Dilemma-style environment.

Main point to land: the incentives are calibrated so that contribution is socially beneficial but individually unattractive.

21. Treatment and Hypotheses

Target time: 5-10 sec

Quick signpost only. I use this slide to move from design mechanics to the treatment structure and empirical predictions.

Main point to land: the next slides connect the design to the hypotheses.

22. Treatments

Target time: 30-40 sec

The treatment is a clean two-by-two design. We vary the domain of disagreement, political versus non-political, and the match type, similar versus opposed. Assignment is based on survey responses, so the manipulation changes who participants believe they are paired with, not the monetary rule itself.

Main point to land: the design varies perceived alignment across both domain and pair type.

23. Hypotheses

Target time: 40-50 sec

H1 predicts more free-riding in similar matches, because aligned partners may expect one another to contribute. H2 predicts that this match effect is stronger in the political domain, where disagreement is more salient. H3 is a broader content effect: contribution may differ between political and non-political movies more generally.

Main point to land: the hypotheses separate match effects, domain effects, and their interaction.

24. Preliminary Results

Target time: 5-10 sec

Quick signpost only. I use this slide to warn the audience that what follows is preliminary and descriptive rather than definitive.

Main point to land: the next slides are pilot-style evidence with clear power limits.

25. Pilot on Prolific: descriptive overview

Target time: 35-45 sec

The completed sample is fifty-three participants, all of whom finished the ten rounds. Average contribution is just over four movies rated out of ten, with a median of four. The table already hints at heterogeneity across cells, but the main message here is modest: there is meaningful behavioral variation, though the sample is still small.

Main point to land: the pilot produces usable variation, but not precise estimates.

26. Results: Hypothesis 1

Target time: 40-50 sec

For H1, contribution is slightly lower in homogeneous than in heterogeneous matches, which goes in the predicted direction, but the estimate is not statistically significant. The power line matters here: to detect a shift of meaningful size, the required sample is much larger than what the pilot provides.

Main point to land: H1 gets directional but weak support, and power is the binding constraint.

27. Results: Hypothesis 2

Target time: 50-60 sec

H2 is the most interesting result. The heterogeneous effect is positive in the non-political condition but negative in the political condition, so the difference across domains is statistically significant, but it runs opposite to the preregistered prediction. That makes this slide interesting, but it also means interpretation matters more than simple hypothesis checking.

Main point to land: the interaction appears real in the pilot, but it points in the opposite direction from the ex ante prediction.

28. Results: Hypothesis 3

Target time: 35-45 sec

For H3, the political-versus-non-political gap is small in both match types, and the interaction is imprecise. The power discussion again matters: even with substantially larger samples, this is the sort of effect that may remain hard to detect cleanly.

Main point to land: there is little evidence for a robust content effect of the H3 type.

29. Interpreting the pilot: mechanism or priming?

Target time: 40-50 sec

This is where I would slow down and show caution. The theory says contributions fall when I expect my partner to contribute more, so similar matches could generate more free-riding. But if monetary incentives are fixed across treatments, then any treatment effect must come from beliefs, interpretation, or framing rather than from a change in the payoff rule itself.

Main point to land: the mechanism is plausibly strategic, but the pilot does not yet isolate that mechanism cleanly.

30. Interpreting the pilot: instruction channel

Target time: 35-45 sec

This slide makes the identification problem explicit. The manipulation may change beliefs about partner behavior and informativeness, but it may also change trust, identity salience, or demand effects because of the wording of the instructions. So the treatment is not purely strategic; it may partly work by priming.

Main point to land: the design may be mixing mechanism and framing, and that needs to be acknowledged directly.

31. Conclusions

Target time: 40-50 sec

The clean summary is this. H1 is directionally consistent with the theory but imprecise. H2 shows a significant difference across domains, but opposite to the preregistered direction. H3 looks small. So the pilot is informative mainly because it identifies where the design is promising, where the power problem is severe, and where interpretation needs to be sharpened before making stronger claims.

Main point to land: the pilot is useful, but the responsible conclusion is cautious and mechanism-focused.

32. Thank you

Target time: 10-15 sec

Thank you very much. I would especially welcome questions on the free-riding mechanism, the treatment interpretation, and how best to separate strategic response from priming in the next iteration.

Main point to land: invite discussion on mechanism and design.

Backup Slides

Backup: Phase 1 ranking screen

Target time: 20-30 sec

I would use this only if someone asks what the ground-truth task looks like in practice. The screenshot lets me show that participants actively construct an ordered preference ranking rather than just ticking boxes.

Main point to land: the accuracy benchmark comes from an explicit ranking task.

Backup: game flow

Target time: 20-30 sec




Figure 1: Ranking screen

This is useful when someone wants the entire experiment in one picture. I would use it to walk through the three phases quickly and show where elicitation, contribution decisions, and payment fit together.

Main point to land: the experiment has a simple sequence despite the theoretical framing.

Backup: matching message

Target time: 20-30 sec

This slide is useful if someone asks how salient the treatment is to participants. It shows that the match context is made explicit by the interface rather than left implicit.

Main point to land: participants are clearly informed about the similar-versus-opposed context.




Figure 2: Matching message

Backup: neutral movie screen

Target time: 20-30 sec

I would use this if someone asks what a non-political decision round looks like.

It helps show that the task remains visually and mechanically similar across domains.

Main point to land: the neutral condition changes content, not the structure of the choice.

Current budget: €2.50
Spent so far: €0.00

In the Shadows

IN THE SHADOWS

MOVIE DESCRIPTION
A violent criminal strikes again, leaving detectives with one night to stop the next attack. With leads collapsing and witnesses too scared to talk, the hunt turns personal. As rain floods the streets and the clock runs out, one investigator follows a clue that could end it or get them killed.

How much would you like to see this movie realized? *

☐ Skip (no cost)

☐ Rate this movie (pay EUR 0.25)

Complete

Figure 3: Neutral movie screen

Backup: political movie screen

Target time: 20-30 sec

This is the parallel slide for the political condition. It is useful if someone wants to see how political content is embedded in the same interface and decision structure.

Main point to land: the political treatment changes the stimulus context while keeping the contribution choice constant.

Backup: click probability

Target time: 25-35 sec

If someone asks for the behavioral meaning of alignment in the model, this is the slide to use. Engagement is highest when the implemented political level is closest to a group's ideal, and it falls as the system moves away from that ideal.

Main point to land: alignment matters in the model because misalignment reduces engagement.

Backup: platform profit

Target time: 25-35 sec

Current budget: €2.25

Spent so far: €0.25

Silicon Circus



MOVIE DESCRIPTION

Silicon Circus loves "global talent," but many workers are immigrants, and some are undocumented. They joke between meetings until panic hits. When an audit arrives, one employee risks everything to protect the people who built the company.

How much would you like to see this movie realized? *

☐ Skip (no cost)
 ☒ Rate this movie (pay EUR 0.25)

If you choose to rate, what is your rating? *

☐ 1 - Strongly dislike
 ☐ 2 - Dislike
 ☐ 3 - Neutral
 ☐ 4 - Like
 ☐ 5 - Strongly like

Complete

Figure 4: Political movie screen

This is the platform side of the same story. The platform responds to user engagement and contributions, but stronger political constraints are costly. So the chosen political level balances revenue incentives against implementation costs.

Main point to land: the platform does not choose ethics in a vacuum; it responds to incentives.

Backup: payoff details

Target time: 20-30 sec

This is useful if someone wants the payoff rule spelled out one symbol at a time. I would use it to show that the notation maps directly onto the incentive explanation already given in the main deck.

Main point to land: the payoff mechanics are transparent and easy to audit.

Timing Guide

Recommended target timings

Slide	Target	Cumulative
1. Title	0:15	0:15
2. Political alignment in recommender systems	0:35	0:50

Slide	Target	Cumulative
3. Political alignment as a public good	0:40	1:30
4. Contribution & Research Questions	0:05	1:35
5. Objective and contribution	0:45	2:20
6. Research questions	0:35	2:55
7. Theoretical Framework	0:05	3:00
8. Model intuition	0:45	3:45
9. A simple model	0:45	4:30
10. Equilibrium intuition	0:50	5:20
11. Main theoretical implication	0:35	5:55
12. Experimental Design	0:05	6:00
13. Experimental design: overview	0:40	6:40
14. Phase 1: Pre-treatment survey & ground truth	0:45	7:25
15. Phase 2: Rating rounds	0:40	8:05
16. Phase 2: The contribution decision	0:35	8:40
17. Phase 2: Movie stimuli	0:35	9:15
18. Phase 3: Individual payoff	0:45	10:00
19. Phase 3: From ratings to accuracy	0:40	10:40
20. Phase 3: Incentive structure	0:40	11:20
21. Treatment and Hypotheses	0:05	11:25
22. Treatments	0:35	12:00
23. Hypotheses	0:45	12:45
24. Preliminary Results	0:05	12:50

Slide	Target	Cumulative
25. Pilot on Prolific: descriptive overview	0:40	13:30
26. Results: Hypothesis 1	0:45	14:15
27. Results: Hypothesis 2	0:55	15:10
28. Results: Hypothesis 3	0:40	15:50
29. Open issue: mechanism or priming?	0:45	16:35
30. Open issue: treatment also works through instructions	0:40	17:15
31. Conclusions	0:45	18:00
32. Thank you	0:10	18:10

This is a realistic full-length version for an 18-minute talk. With natural pauses and one or two audience-facing transitions, it lands around 18-19 minutes.

Rough delivery modes

Mode	Practical use
Compressed 15 min	Shorten slides 10, 18-20, 27, 29, and 30, and treat the section-divider slides as near-silent transitions.
Standard 18 min	Use the timings above and give the most space to slides 10, 18-20, 27, and 29.
Relaxed 20 min	Add a few extra sentences on the theoretical mechanism, the payoff logic, and the interpretation caveat in slides 29-30.

If time runs long

Shorten first: - slide 10, Equilibrium intuition - slide 18, Phase 3: Individual payoff - slide 27, Results: Hypothesis 2 - slide 30, Open issue: treatment also works through instructions

Skip first under real pressure: - slide 17, Phase 2: Movie stimuli

Most useful backup slides: - Backup: click probability - Backup: platform profit
- Backup: payoff details